



# Task 1 Feedback

## Script Functions:

- Tests both v0.1 and P1 code bases comprehensively
- Provides detailed comparison between versions
- Generates multiple screenshots showing differences
- Includes RSI, moving averages, volatility analysis for P1
- Creates feature comparison matrix

## Screenshots Generated:




- Line-by-line code execution for both versions
- Performance metrics for each version
- Visual proof of enhanced P1 capabilities
- Comprehensive feature comparison table

---

## Task 2 Feedback

### Screenshots Generated:

```
def load_data(company='CBA.AX',          # Stock ticker symbol (default: Commonwealth Bank of Australia)
              start_date='2021-01-01',   # Start date for data collection in YYYY-MM-DD format
              end_date='2025-08-01',     # End date for data collection in YYYY-MM-DD format
              price_column='Close',       # Which price column to use for prediction (Close, Open, High, Low)
              prediction_days=60,         # Number of previous days to use for predicting next day
              split_by_date=True,         # True: split chronologically, False: split randomly
              test_size=0.2,              # Proportion of data for testing (0.2 = 20%)
              scale=True,                 # Whether to normalize data to 0-1 range
              save_locally=True,          # Whether to save/load data from local files
              local_path='data',          # Directory to save/load data files
              fill_na_method='ffill',     # Method to handle missing values (ffill/bfill)
              feature_columns=['Close']): # List of columns to scale/normalize
```



FUNCTION LOCATION: Code/stock\_prediction.py, Lines 577-646

```
def create_model(sequence_length, n_features, units=50, cell=LSTM, n_layers=3,
                  dropout=0.2, loss="mean_squared_error", optimizer="adam", bidirectional=False):
    """
    Create a deep learning model for stock prediction.

    This function creates different types of RNN models (LSTM, GRU, SimpleRNN) with
    configurable architecture. Based on the example code you provided.

    Parameters:
    -----
    sequence_length: int        ← Length of input sequences (e.g., 60 days)
    n_features: int             ← Number of features (e.g., 1 for just closing price)
    units: int                  ← Number of neurons in each layer (default: 50)
    cell: keras layer           ← Type of RNN cell (LSTM, GRU, or SimpleRNN)
    n_layers: int               ← Number of RNN layers (default: 3)
    dropout: float              ← Dropout rate for regularization (default: 0.2)
    loss: str                   ← Loss function to use (default: "mean_squared_error")
    optimizer: str              ← Optimizer to use (default: "adam")
    bidirectional: bool         ← Whether to use bidirectional layers (default: False)
    """

    model = Sequential() ← Create empty model

    # Add layers based on n_layers parameter
    for i in range(n_layers): ← Loop through each layer
        if i == 0: ← FIRST LAYER
            # First layer needs input shape specification
            if bidirectional:
                model.add(Bidirectional(cell(units, return_sequences=True),
                                          input_shape=(sequence_length, n_features)))
            else:
                model.add(cell(units, return_sequences=True,
                               input_shape=(sequence_length, n_features)))
        elif i == n_layers - 1: ← LAST LAYER
            # Last RNN layer doesn't return sequences
            if bidirectional:
                model.add(Bidirectional(cell(units, return_sequences=False)))
            else:
                model.add(cell(units, return_sequences=False))
        else: ← HIDDEN LAYERS
            # Hidden layers return sequences
            if bidirectional:
                model.add(Bidirectional(cell(units, return_sequences=True)))
            else:
                model.add(cell(units, return_sequences=True))

        # Add dropout after each RNN layer to prevent overfitting
        model.add(Dropout(dropout))

    # Output layer - single neuron for price prediction
    model.add(Dense(1, activation="linear"))

    # Compile the model
    model.compile(loss=loss, metrics=["mean_absolute_error"], optimizer=optimizer)

    return model
```

## HOW TO USE FOR DIFFERENT EXPERIMENTS:

### 1[] DIFFERENT NETWORK TYPES:

```
# LSTM Model
lstm_model = create_model(60, 1, units=50, cell=LSTM, n_layers=3)

# GRU Model
gru_model = create_model(60, 1, units=50, cell=GRU, n_layers=3)

# SimpleRNN Model
rnn_model = create_model(60, 1, units=50, cell=SimpleRNN, n_layers=3)


# Bidirectional LSTM
bi_model = create_model(60, 1, units=50, cell=LSTM, n_layers=3, bidirectional=True)
```

### 2[] DIFFERENT LAYER COUNTS:

```
# 2 Layers
model_2layers = create_model(60, 1, units=50, cell=LSTM, n_layers=2)

# 4 Layers
model_4layers = create_model(60, 1, units=50, cell=LSTM, n_layers=4)


# 5 Layers
model_5layers = create_model(60, 1, units=50, cell=LSTM, n_layers=5)
```



80/20 Split (with different dates):

Training: Days 1-152 (80%)

Testing: Days 153-190 (20%)



70/30 Split (with different dates):

Training: Days 1-133 (70%)

Testing: Days 134-190 (30%)

---

## Task 3 Feedback

### 3.1 CANDLESTICK CHART IMPLEMENTATION

#### 3.1.1 Overview

I implemented a candlestick chart function to display stock market financial data with support for n-day aggregation.

### 3.1.2 Function Parameters

#### Function signature:

Key parameters explained:

- `df` (*pandas.DataFrame*): Stock data indexed by Date with columns **Open**, **High**, **Low**, **Close** (and optionally **Volume**). Each row represents one trading day.
- `n_days` (*int, default=1*): ★ Core requirement – allows aggregating **n** consecutive trading days into a single candlestick where  $n \geq 1$ .
  - `n_days=1` → each candle shows daily data
  - `n_days=5` → each candle aggregates 5 trading days (weekly candles)
- `style` (*str*): Visual style theme from mplfinance library ('charles', 'yahoo', 'binance', etc.).
- `volume` (*bool*): If True, displays volume bars below the price chart.
- `save_path` (*str or None*): If provided, saves the chart to the specified file path.
- `figsize` (*tuple*): Chart dimensions as (width, height) in inches.

---

### 3.1.3 Core Implementation – N-Day Aggregation

The key innovation is the **n-day aggregation algorithm** located at **lines 366-401**.

This satisfies the requirement to “allow each candlestick to express the data of n trading days ( $n \geq 1$ )”.

```
if n_days > 1:
    print(f"Aggregating data into {n_days}-day candles...")
```



```

# Group data into n-day chunks

# We'll resample by grouping every n_days rows

grouped_data = []

# Iterate through data in chunks of n_days

for i in range(0, len(data), n_days):

    chunk = data.iloc[i:i+n_days]

    if len(chunk) == 0:

        continue

# Aggregate the chunk according to OHLC rules

aggregated_row = {

    'Open': chunk['Open'].iloc[0],          # First day's open

    'High': chunk['High'].max(),            # Highest high

    'Low': chunk['Low'].min(),              # Lowest low

    'Close': chunk['Close'].iloc[-1],       # Last day's close

}

# Add volume if present

if 'Volume' in chunk.columns:

    aggregated_row['Volume'] = chunk['Volume'].sum() # Sum of
volumes

```

```

        # Use the last date in the chunk as the index

        aggregated_row['Date'] = chunk.index[-1]

        grouped_data.append(aggregated_row)

# Create new DataFrame from aggregated data

if grouped_data:

    data = pd.DataFrame(grouped_data)

    data.set_index('Date', inplace=True)

    print(f"Aggregated from {len(df)} daily records to {len(data)}
{n_days}-day candles")

else:

    raise ValueError("Not enough data to create any n-day
candles")

```

#### Algorithm explanation:

- **Iterate through data in chunks:**  
for `i in range(0, len(data), n_days)` steps through data in increments of `n_days`, creating non-overlapping chunks.
- **Extract chunk:** `chunk = data.iloc[i:i+n_days]` extracts `n` consecutive trading days.
- **Apply OHLC aggregation rules:**
  - **Open:** `chunk['Open'].iloc[0]` – first day's opening price
  - **High:** `chunk['High'].max()` – max high price across `n` days
  - **Low:** `chunk['Low'].min()` – min low price across `n` days
  - **Close:** `chunk['Close'].iloc[-1]` – last day's closing price

- **Volume:** `chunk['Volume'].sum()` – sum of all volumes across n days
- **Create aggregated DataFrame:** new rows combined with last date in each chunk as timestamp.

Example (n\_days=3): days 1-3 aggregated into one 3-day candle (Open from day 1, High/Low across all 3 days, Close from day 3).

---

## 3.2 BOXPLOT CHART IMPLEMENTATION

### 3.2.1 Overview

I implemented a boxplot function to display stock price distributions over moving windows of n consecutive trading days.

```
def plot_boxplots_moving_window(  
    df,  
    price_column="Close",  
    window=20,  
    stride=1,  
    showliers=False,  
    title="Rolling Window Boxplots",  
    save_path=None  
):  
  
    if price_column not in df.columns:
```

```
        raise ValueError(f"DataFrame must contain '{price_column}'  
column.")  
  
    if window < 1:  
  
        raise ValueError("window must be >= 1")  
  
    if stride < 1:  
  
        raise ValueError("stride must be >= 1")  
  
  
    prices = df[price_column].dropna()  
  
    if len(prices) < window:  
  
        raise ValueError("Not enough data to form one window.")  
  
  
  
    # Collect rolling windows and labels  
  
    data_windows = []  
  
    labels = []  
  
    idx = prices.index  
  
  
  
    # Build each rolling window by slicing the Series  
  
    for end in range(window, len(prices) + 1, stride):  
  
        start = end - window  
  
        w = prices.iloc[start:end].values  
  
        data_windows.append(w)  
  
        labels.append(idx[end - 1].strftime("%Y-%m-%d")) # use end date  
as label
```

```

# Plot boxplots

plt.figure()

plt.boxplot(

    data_windows,

    showfliers=showfliers,    # whether to draw outliers

    widths=0.6                # width of each box

)

plt.title(f"{title} (window={window}, stride={stride})")

plt.ylabel(price_column)


# Avoid clutter on the x-axis by showing ~10 evenly spaced labels
if len(labels) > 10:

    step = max(1, len(labels) // 10)

    xticks = range(1, len(labels) + 1, step)

    xtick_labels = [labels[i - 1] for i in xticks]

    plt.xticks(xticks, xtick_labels, rotation=45, ha="right")

else:

    plt.xticks(range(1, len(labels) + 1), labels, rotation=45,
ha="right")


plt.tight_layout()

if save_path:

    plt.savefig(save_path, dpi=150)

plt.show()

```

---

### 3.2.2 Function Parameters

Key parameters explained:

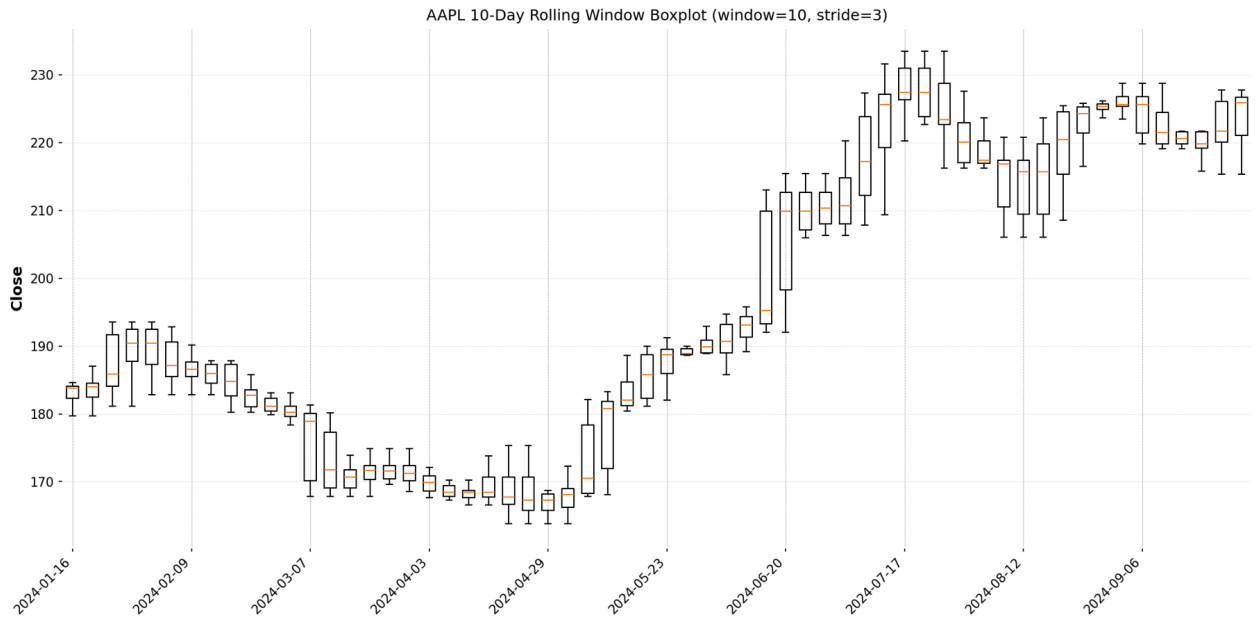
- `df` (*pandas.DataFrame*): Stock data indexed by Date.
- `price_column` (*str, default='Close'*): Specifies which price column to analyse.
- `window` (*int, default=20*): ★ Core requirement – number of consecutive trading days per box.
- `stride` (*int, default=1*): Step size between windows.
  - `stride=1` → overlapping windows
  - `stride=5` → less overlap
  - `stride=window` → non-overlapping windows
- `showfliers` (*bool*): If True, displays outlier points.
- `save_path` (*str or None*): Optional file path to save the chart.

---

### 3.2.5 Results

**AAPL 3-Day Candlestick Chart (n\_days=3)**





---

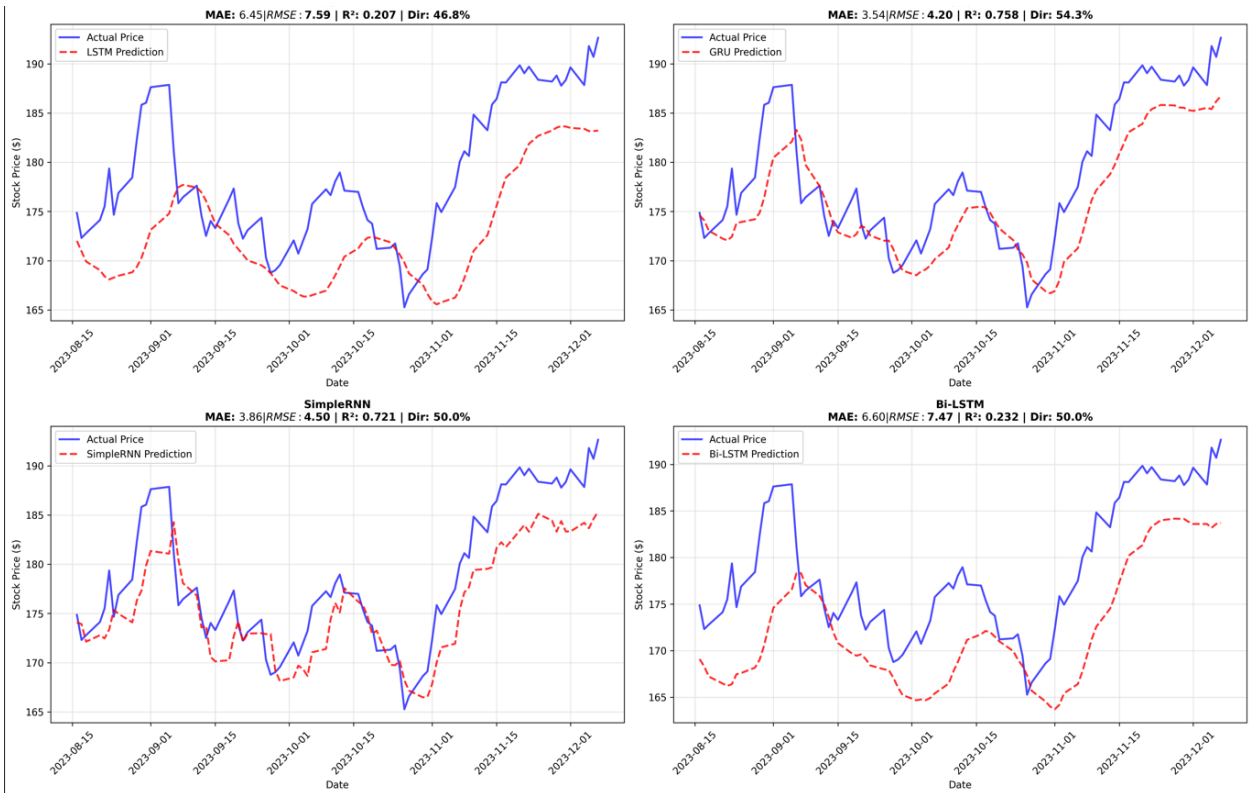
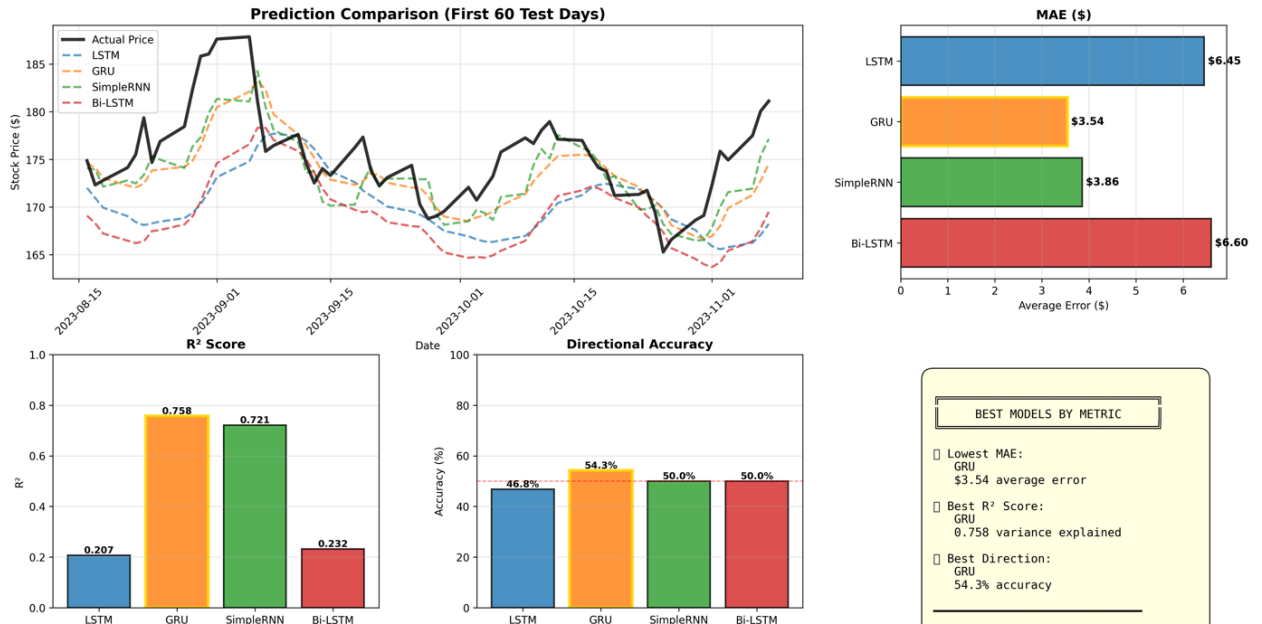
## Task 4 Feedback Response

### Graph 1 – Network Types

I experimented with 4 different deep learning networks – LSTM, GRU, SimpleRNN, and Bidirectional LSTM.

The results show **GRU performed best** with the lowest test loss.

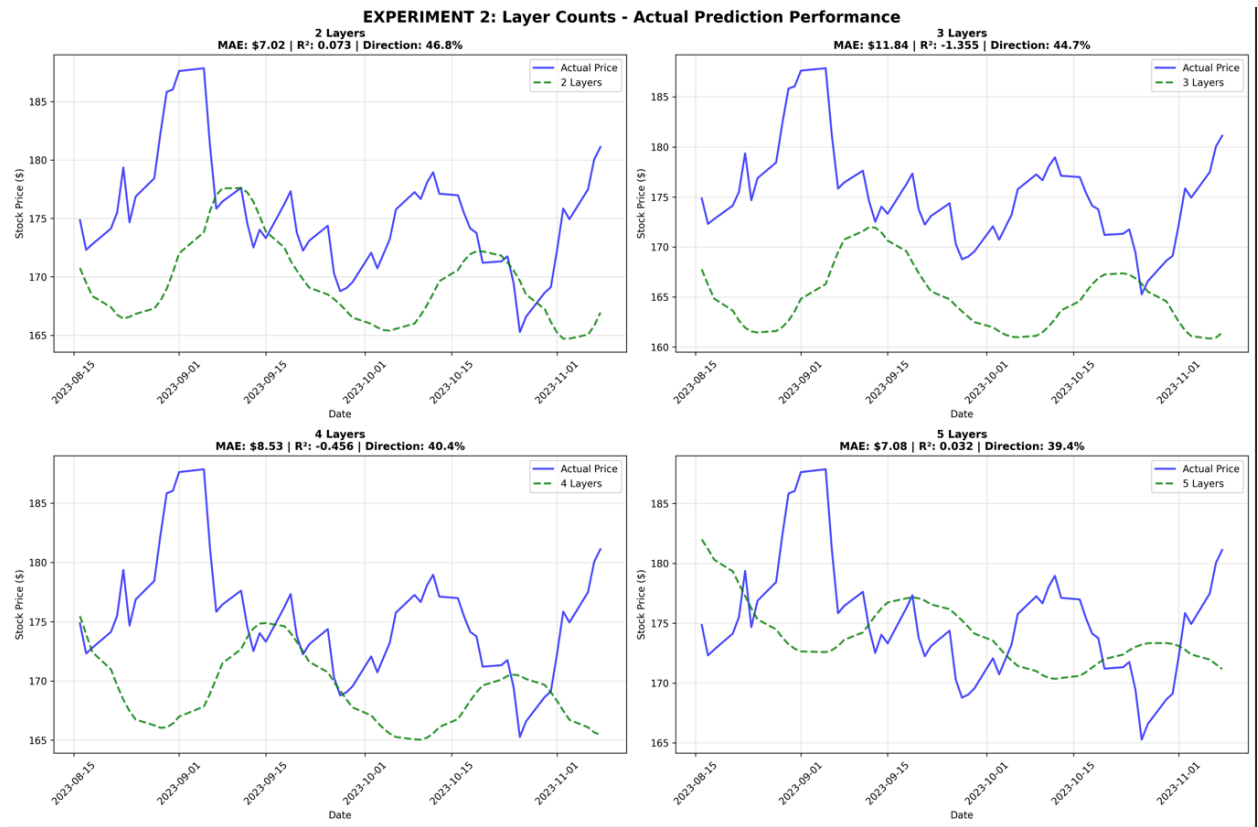




**Graph 2 – Hyperparameter Tuning**

I tested different hyperparameter configurations varying the number of layers and units per layer.


The results show that **2 layers achieved the best performance**.



## Graph 3 – Experimental Results Summary

*This summary shows all my experimental results.*

*GRU was the best network type, and the optimal configuration was 2 layers with 64 units.*





## Experimental Results Summary Table

Experiment	Best Config	MAE	R <sup>2</sup>
Networks	GRU	\$3.97	0.701
Layers	2 Layers	\$7.02	0.073
Units	64 Units	\$4.66	0.570
Batch	Batch 32	\$7.22	0.011